

SATHVIK KONDA

AI/ML Engineer – GenAI, RAG & Agentic Systems
+91-93925-29828 sathvik8317@gmail.com Hyderabad, Telangana
LinkedIn: linkedin.com/in/sathvik-konda-b890121a1 GitHub: sathvik8317
HackerRank: sathvik8317 LeetCode: sathvik8317 X/Threads: @sathvik8317

PROFILE

B.Tech Computer Science (Data Science) graduate working as a full-time Software Engineer while building a focused portfolio in applied AI/ML: retrieval-augmented generation, agentic systems, the Model Context Protocol (MCP), LLM fine-tuning, and computer vision, with several projects built and evaluated on constrained consumer hardware. Publishes build-in-public technical writing on LLM infrastructure and agentic AI on Medium, X, and Threads.

EDUCATION

- VNR Vignana Jyothi Institute of Engineering and Technology**, Hyderabad *Nov 2022 – Jun 2026*
B.Tech – Computer Science and Engineering (Data Science); CGPA: 8.61

EXPERIENCE

- Associate Developer – Issquared**, Hyderabad *Jul 2026 – Present*
Building and maintaining production applications on a .NET / ASP.NET Core and Angular stack; designed reference APIs (layered architecture, EF Core, DTOs, repository pattern). Applying MLOps and agentic-tooling practices from personal research to internal engineering workflows.

SKILLS

- AI/ML & GenAI:** LLMs, RAG, Agentic AI, Model Context Protocol (MCP), LoRA/QLoRA fine-tuning (Unsloth, TRL), Prompt Engineering, Computer Vision, MLOps
- Languages & Frameworks:** Python, C#, SQL, TypeScript, .NET / ASP.NET Core, Angular, PyTorch, TensorFlow/Keras, Scikit-learn, LangChain, FastAPI
- AI Infra & Tools:** ChromaDB, LanceDB, Ollama, FastMCP, Docker, Git, Streamlit, Power BI
- Soft Skills:** Teamwork, Problem Solving, Time Management, Strategic Thinking

ACHIEVEMENTS

- HackerRank:** 4* in Problem Solving; 5* in SQL (Basic, Intermediate & Advanced)
- Presented research paper "*A Modular Time-Aware SMOTE and Autoencoder-XGBoost Framework for Interpretable Credit Card Fraud Detection*" at **ICACECS Conference 2025**.
- Build-in-public technical writer: publishes AI/ML infrastructure, local-first inference, and agentic-systems content on Medium, X, and Threads (@sathvik8317).

PROJECTS

- Multimodal Search Engine** – GitHub
Hybrid document search using Cohere Embed v4 and LanceDB, combining reciprocal rank fusion (RRF) with Cohere reranking and tree-sitter AST-based code chunking; 183 passing tests.
- Multimodal Document Intelligence Agent**
Vision-based invoice/receipt extraction pipeline using Qwen3-VL 2B (Ollama) as the local backend, benchmarked against SROIE and CORD ground truth with a Streamlit evaluation dashboard.
- Local-First Inference Proxy**
FastAPI service that escalates from a local Qwen3-VL model to the OpenAI API using logprob-based and heuristic confidence scoring; demonstrates an MLOps cost/accuracy tradeoff pattern on a 4GB VRAM laptop GPU.
- DesiTutor – Fine-Tuned Hinglish AI Tutor**
QLoRA fine-tune of Qwen3-4B-Instruct (Unsloth + TRL) with LLM-as-judge evaluation and positional-bias correction; exported to GGUF for local Ollama deployment.
- shared-brain-mcp**
FastMCP stdio MCP server backed by ChromaDB for persistent, cross-session engineering memory, with Markdown/YAML as the source of truth.

CERTIFICATIONS

- NPTEL Certification: Data Science For Engineers
- Coursera: Machine Learning by Andrew Ng
- Udemy: Microsoft Certified Data Analyst Associate (Power BI)

INTERESTS

LLM infrastructure and inference internals; local-first and constrained-hardware AI; agentic systems; Chess; Photography